41. Pass4for

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 68 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 68
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer is working with two tables. Each of these tables is displayed below in its entirety.

sales

customer_id	spend	units
a1	28.94	7
a3	874.12	23
a4	8.99	1

favorite_stores

customer_id	store_id
a1	s1
a2	s1
a4	s2

The data engineer runs the following query to join these tables together:

```
SELECT
    sales.customer_id,
    sales.spend,
    favorite_stores.store_id
FROM sales
LEFT JOIN favorite_stores
ON sales.customer_id = favorite_stores.customer_id;
```

Which of the following will be returned by the above query?

- A.

customer_id	spend	store_id
a1	28.94	s1
a4	8.99	s2

- B.

customer_id	spend	units	store_id
a1	28.94	7	s1
a4	8.99	1	s2

- C.

customer_id	spend	store_id
a1	28.94	s1
a3	874.12	NULL
a4	8.99	s2

- D.

customer_id	spend	store_id
a1	28.94	s1
a2	NULL	s1
a3	874.12	NULL
a4	8.99	s2

- E.

customer_id	spend	store_id
a1	28.94	s1
a2	NULL	s1
a4	8.99	s2

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 189 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 189
Topic #: 1
[All Certified Data Engineer Associate Questions]

What is the structure of an Asset Bundle?

A. A Docker image containing runtime environments and the source code of the assets

B. A compressed archive (ZIP) that solely contains workspace assets without any accompanying metadata

C. A single plain text file enumerating the names of assets to be migrated to a new workspace

D. A YAML configuration file that specifies the artifacts, resources, and configurations for the project

Show Suggested Answer

by taimurkhan222 at Nov. 10, 2025, 3:13 a.m.

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 193 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 193
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineering project involves processing large batches of data on a daily schedule using ETL. The jobs are resource-intensive and vary in size, requiring a scalable, cost-efficient compute solution that can automatically scale based on the workload.

Which compute approach will satisfy the needs described?

A. Job Cluster
B. Dedicated Cluster
C. All-Purpose Cluster
D. Databricks SQL Serverless

Show Suggested Answer

by taimurkhan222 at *Nov. 10, 2025, 3:14 a.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 194 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 194
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer is configuring Delta Sharing for a multi-team project where teams from different departments will need to access shared data. The data engineer has successfully created a Unity Catalog metastore and is now setting up the Delta Share. The goal is to ensure that internal teams can access the shared data with full permissions, while external partners can only read the data.

Which action should the Data Engineer take to configure the sharing correctly?

A. Create a Delta Share, set up a secure access URL for internal teams and external partners, and distribute the URL to provide them access to the shared data.
B. Create a Delta Share, add the internal team's tables and views, and assign READ/WRITE permissions to both external partners and internal teams.
C. Assign READ permissions to external partners through the Delta Share and READ/WRITE permissions to internal teams, and ensure the correct tables and views are shared.
D. Grant READ permissions to external partners and READ/WRITE permissions to internal teams through the Delta Share.

Show Suggested Answer

by taimurkhan222 at *Nov. 10, 2025, 3:13 a.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 186 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 186
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer has inherited a Databricks pipeline from a previous team. The pipeline is missing its SLAs, and an initial investigation has identified memory spills in Spark. These increased runtimes are also driving up costs. The data engineer needs to reduce the runtime without significantly increasing costs.

What should the data engineer do first to address this issue?

A. Tweak the "spark.sql.shuffle.partitions" configuration.
B. Enable autoscaling in the cluster to match the requirements
C. Use a photon enabled execution engine.
D. Ensure the cluster uses memory optimized node types.

Show Suggested Answer

by taimurkhan222 at *Nov. 10, 2025, 3:14 a.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 185 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 185
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer is working in a Python notebook on Databricks to process data, but notices that the output is not as expected. The data engineer wants to investigate the issue by stepping through the code and checking the values of certain variables during execution.

Which tool should the data engineer use to inspect the code execution and variables in real-time?

A. Cluster Logs
B. Job Execution Dashboard
C. Python Notebook Interactive Debugger
D. SQL Analytics

Show Suggested Answer

by taimurkhan222 at *Nov. 10, 2025, 3:14 a.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 183 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 183
Topic #: 1
[All Certified Data Engineer Associate Questions]

What is the primary function of the Silver layer in the Databricks medallion architecture?

A. Store historical data solely for auditing purposes
B. Aggregate and enrich data for business analytics
C. Validate, clean, and deduplicate data for further processing
D. Ingest raw data in its original state

Show Suggested Answer

by taimurkhan222 at *Nov. 10, 2025, 3:13 a.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 182 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 182
Topic #: 1
[All Certified Data Engineer Associate Questions]

A Data Engineer is designing Bronze layer in Databricks Medallion Architecture. The raw data is collected from multiple sources (Clickstream in JSON, Transactions in CSV).

The task is to design the Bronze layer of the Medallion Architecture to ingest and store this raw data for further processing.

Which operation applies to the Bronze layer?

A. Ingest raw data without transformations, preserving original schemas, and store in Delta format.
B. Clean and standardize raw data by removing null values and enforcing schemas.
C. Apply complex business logic to enrich raw data with customer segmentation labels.
D. Aggregate and transform source data to calculate daily sales performance metrics.

Show Suggested Answer

by taimurkhan222 at *Nov. 10, 2025, 3:14 a.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 176 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 176
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer needs to provide access to a group named manufacturing-team. The team needs privileges to create tables in the quality schema.

Which set of SQL commands will grant a group named manufacturing-team to create tables in a schema named production with the parent catalog named manufacturing with the least privileges?

A. GRANT CREATE TABLE ON SCHEMA manufacturing.quality TO manufacturing-team; GRANT USE SCHEMA ON SCHEMA manufacturing.quality TO manufacturing-team; GRANT USE CATALOG ON CATALOG manufacturing TO manufacturing-team;

B. GRANT USE TABLE ON SCHEMA manufacturing.quality TO manufacturing-team; GRANT USE SCHEMA ON SCHEMA manufacturing.quality TO manufacturing-team; GRANT USE CATALOG ON CATALOG manufacturing TO manufacturing-team;

C. GRANT CREATE TABLE ON SCHEMA manufacturing.quality TO manufacturing-team; GRANT CREATE SCHEMA ON SCHEMA manufacturing.quality TO manufacturing-team; GRANT CREATE CATALOG ON CATALOG manufacturing TO manufacturing-team;

D. GRANT CREATE TABLE ON SCHEMA manufacturing.quality TO manufacturing-team; GRANT CREATE SCHEMA ON SCHEMA manufacturing.quality TO manufacturing-team; GRANT USE CATALOG ON CATALOG manufacturing TO manufacturing-team;

Show Suggested Answer

by [deleted] at *Oct. 4, 2025, 9:02 a.m.*

Q173

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 1
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer is working with two tables. Each of these tables is displayed below in its entirety.

sales

customer_id	spend	units
a1	28.94	7
a3	874.12	23
a4	8.99	1

favorite_stores

customer_id	store_id
a1	s1
a2	s1
a4	s2

The data engineer runs the following query to join these tables together:

```
SELECT
    sales.customer_id,
    sales.spend,
    favorite_stores.store_id
FROM sales
LEFT JOIN favorite_stores
ON sales.customer_id = favorite_stores.customer_id;
```

Which of the following will be returned by the above query?

A.

customer_id	spend	store_id
a1	28.94	s1
a4	8.99	s2

B.

customer_id	spend	store_id
a1	28.94	s1
a2	NULL	s1
a4	8.99	s2

C.

customer_id	spend	store_id
a1	28.94	s1
a3	874.12	NULL
a4	8.99	s2

D.

customer_id	spend	store_id
a1	28.94	s1
a2	NULL	s1
a3	874.12	NULL
a4	8.99	s2

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 169 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 10
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer wants to create a relational object by pulling data from two tables. The relational object does not need to be used by other data engineers in other sessions. In order to save on storage costs, the data engineer wants to avoid copying and storing physical data.

Which of the following relational objects should the data engineer create?

A. Spark SQL Table
B. View
C. Delta Table
D. Temporary view

Show Suggested Answer

by MultiCloudIronMan at *Dec. 16, 2024, 4:06 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 168 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 15
Topic #: 1
[All Certified Data Engineer Associate Questions]

What is stored in the Databricks customer's cloud account?

A. Databricks web application

B. Cluster management metadata

C. Notebooks

D. Data

Show Suggested Answer

by MultiCloudIronMan at *Dec. 16, 2024, 4:05 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 167 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 17
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer is designing a data pipeline. The source system generates files in a shared directory that is also used by other processes. As a result, the files should be kept as is and will accumulate in the directory. The data engineer needs to identify which files are new since the previous run in the pipeline, and set up the pipeline to only ingest those new files with each run.

Which of the following tools can the data engineer use to solve this problem?

A. Unity Catalog

B. Delta Lake

C. Databricks SQL

D. Auto Loader

Show Suggested Answer

by MultiCloudIronMan at *Dec. 16, 2024, 4:04 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 150 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 53
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer needs to apply custom logic to string column city in table stores for a specific use case. In order to apply this custom logic at scale, the data engineer wants to create a SQL user-defined function (UDF).

Which of the following code blocks creates this SQL UDF?

A.
```
CREATE FUNCTION combine_nyc(city STRING)
RETURNS STRING
RETURN CASE
  WHEN city = "brooklyn" THEN "new york"
  ELSE city
END;
```

B.
```
CREATE UDF combine_nyc(city STRING)
RETURNS STRING
CASE
  WHEN city = "brooklyn" THEN "new york"
  ELSE city
END;
```

C.
```
CREATE FUNCTION combine_nyc(city STRING)
RETURN CASE
  WHEN city = "brooklyn" THEN "new york"
  ELSE city
END;
```

D.
```
CREATE UDF combine_nyc(city STRING)
RETURNS STRING
RETURN CASE
  WHEN city = "brooklyn" THEN "new york"
  ELSE city
END;
```

Show Suggested Answer

by Manish_Kum at *Dec. 5, 2024, 11:22 p.m.*