EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 100 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 100
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer has been given a new record of data:

id STRING = 'a1'
rank INTEGER = 6
rating FLOAT = 9.4

Which SQL commands can be used to append the new record to an existing Delta table my_table?

A. INSERT INTO my_table VALUES ('a1', 6, 9.4)

B. INSERT VALUES ('a1', 6, 9.4) INTO my_table

C. UPDATE my_table VALUES ('a1', 6, 9.4)

D. UPDATE VALUES ('a1', 6, 9.4) my_table

Show Suggested Answer

by MDWPartners at *May 25, 2024, 8:13 p.m.*

Comments

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 101 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 101
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer has realized that the data files associated with a Delta table are incredibly small. They want to compact the small files to form larger files to improve performance.

Which keyword can be used to compact the small files?

A. OPTIMIZE

B. VACUUM

C. COMPACTION

D. REPARTITION

Show Suggested Answer

by MDWPartners at *May 25, 2024, 8:13 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 102 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 102
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer wants to create a data entity from a couple of tables. The data entity must be used by other data engineers in other sessions. It also must be saved to a physical location.

Which of the following data entities should the data engineer create?

A. Table

B. Function

C. View

D. Temporary view

Show Suggested Answer

by Dip1994 at *June 25, 2024, 3:10 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 103 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 103
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer runs a statement every day to copy the previous day's sales into the table transactions. Each day's sales are in their own file in the location "/transactions/raw".

Today, the data engineer runs the following command to complete this task:

```
COPY INTO transactions
FROM "/transactions/raw"
FILEFORMAT = PARQUET;
```

After running the command today, the data engineer notices that the number of records in table transactions has not changed.

What explains why the statement might not have copied any new records into the table?

A. The format of the files to be copied were not included with the FORMAT_OPTIONS keyword.

B. The COPY INTO statement requires the table to be refreshed to view the copied rows.

C. The previous day's file has already been copied into the table.

D. The PARQUET file format does not support COPY INTO.

Show Suggested Answer

by MDWPartners at *May 25, 2024, 8:14 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 104 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 104
Topic #: 1
[All Certified Data Engineer Associate Questions]

Which command can be used to write data into a Delta table while avoiding the writing of duplicate records?

A. DROP

B. INSERT

C. MERGE

D. APPEND

Show Suggested Answer

by MDWPartners at *May 25, 2024, 8:13 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 105 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 105
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data analyst has created a Delta table sales that is used by the entire data analysis team. They want help from the data engineering team to implement a series of tests to ensure the data is clean. However, the data engineering team uses Python for its tests rather than SQL.

Which command could the data engineering team use to access sales in PySpark?

A. SELECT * FROM sales

B. spark.table("sales")

C. spark.sql("sales")

D. spark.delta.table("sales")

Show Suggested Answer

by MDWPartners at *May 25, 2024, 8:14 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 106 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 106
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer has created a new database using the following command:

CREATE DATABASE IF NOT EXISTS customer360;

In which location will the customer360 database be located?

A. dbfs:/user/hive/database/customer360

B. dbfs:/user/hive/warehouse

C. dbfs:/user/hive/customer360

D. dbfs:/user/hive/database

Show Suggested Answer

by MDWPartners at *May 25, 2024, 8:14 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 107 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 107
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer is attempting to drop a Spark SQL table my_table and runs the following command:

DROP TABLE IF EXISTS my_table;

After running this command, the engineer notices that the data files and metadata files have been deleted from the file system.

What is the reason behind the deletion of all these files?

A. The table was managed

B. The table's data was smaller than 10 GB

C. The table did not have a location

D. The table was external

Show Suggested Answer

by jetplanes at *May 9, 2024, 6:19 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 108 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 108
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer needs to create a table in Databricks using data from a CSV file at location /path/to/csv.

They run the following command:

```
CREATE TABLE new_table
_____
OPTIONS (
  header = "true",
  delimiter = "|"
)
LOCATION "path/to/csv"
```

Which of the following lines of code fills in the above blank to successfully complete the task?

A. FROM "path/to/csv"

B. USING CSV

C. FROM CSV

D. USING DELTA

Show Suggested Answer

by CommanderBigMac at *Sept. 12, 2024, 5:36 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 109 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 109
Topic #: 1
[All Certified Data Engineer Associate Questions]

What is a benefit of creating an external table from Parquet rather than CSV when using a CREATE TABLE AS SELECT statement?

A. Parquet files can be partitioned

B. Parquet files will become Delta tables

C. Parquet files have a well-defined schema

D. Parquet files have the ability to be optimized

Show Suggested Answer

by mgari at *May 14, 2024, 2:12 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 110 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 110
Topic #: 1
[All Certified Data Engineer Associate Questions]

Which SQL keyword can be used to convert a table from a long format to a wide format?

A. TRANSFORM

B. PIVOT

C. SUM

D. CONVERT

Show Suggested Answer

by 9d4d68a at *Aug. 26, 2024, 10:30 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 111 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 111
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer has a Python variable table_name that they would like to use in a SQL query. They want to construct a Python code block that will run the query using table_name.

They have the following incomplete code block:

____(f"SELECT customer_id, spend FROM {table_name}")

What can be used to fill in the blank to successfully complete the task?

A. spark.delta.sql

B. spark.sql

C. spark.table

D. dbutils.sql

Show Suggested Answer

by Stefan94 at *Sept. 20, 2024, 3:46 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 112 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 112
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer is working with two tables. Each of these tables is displayed below in its entirety.

```
sales
customer_idspend units
a1          28.94 7
a3          874.1223
a4          8.99  1

favorite_stores


customer_idstore_id
a1          s1
a2          s1
a4          s2
```

The data engineer runs the following query to join these tables together:

```
SELECT
    sales.customer_id,
    sales.spend,
    favorite_stores.store_id
FROM sales
LEFT JOIN favorite_stores
ON sales.customer_id = favorite_stores.customer_id;
```

A.
```
customer_idspendstore_id
a1          28.94s1
a2          NULL s1
a4          8.99 s2
```

B.
```
customer_idspendstore_id
a1          28.94 s1
a4          8.99  s2
```

C.
```
customer_idspend store_id
a1          28.94 s1
a3          874.12NULL
a4          8.99  s2
```

D.
```
customer_idspend store_id
a1          28.94 s1
a2          NULL  s1
a3          874.12NULL
a4          8.99  s2
```

Show Suggested Answer

by b79962e at *May 22, 2024, 10:07 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 113 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 113
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer needs to apply custom logic to identify employees with more than 5 years of experience in array column employees in table stores. The custom logic should create a new column exp_employees that is an array of all of the employees with more than 5 years of experience for each row. In order to apply this custom logic at scale, the data engineer wants to use the FILTER higher-order function.

Which code block successfully completes this task?

A.
```
SELECT
    store_id,
    employees,
    FILTER (employees, i -> i.years_exp > 5) AS exp_employees
FROM stores;
```

B.
```
SELECT
    store_id,
    employees,
    FILTER (exp_employees, i -> i.years_exp > 5) AS exp_employees
FROM stores;
```

C.
```
SELECT
    store_id,
    employees,
    FILTER (employees, years_exp > 5) AS exp_employees
FROM stores;
```

D.
```
SELECT
    store_id,
    employees,
    CASE WHEN employees.years_exp > 5 THEN employees
         ELSE NULL
    END AS exp_employees
FROM stores;
```

Show Suggested Answer

by CommanderBigMac at *Sept. 12, 2024, 5:50 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 114 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 114
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer that is new to using Python needs to create a Python function to add two integers together and return the sum?

Which code block can the data engineer use to complete this task?

A.
```
function add_integers(x, y):
    return x + y
```

B.
```
def add_integers(x, y):
    print(x + y)
```

C.
```
def add_integers(x, y):
    x + y
```

D.
```
def add_integers(x, y):
    return x + y
```

Show Suggested Answer

by jetplanes at *May 9, 2024, 6:21 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 115 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 115
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer has configured a Structured Streaming job to read from a table, manipulate the data, and then perform a streaming write into a new table.

The code block used by the data engineer is below:

```
(spark.table("sales")
    .withColumn("avg_price", col("sales") / col("units"))
    .writeStream
    .option("checkpointLocation", checkpointPath)
    .outputMode("complete")
    .______
    .table("new_sales")
)
```

Which line of code should the data engineer use to fill in the blank if the data engineer only wants the query to execute a micro-batch to process data every 5 seconds?

A. trigger("5 seconds")

B. trigger(continuous="5 seconds")

C. trigger(once="5 seconds")

D. trigger(processingTime="5 seconds")

Show Suggested Answer

by 9d4d68a at *Aug. 27, 2024, 12:45 a.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 116 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 116
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer is maintaining a data pipeline. Upon data ingestion, the data engineer notices that the source data is starting to have a lower level of quality. The data engineer would like to automate the process of monitoring the quality level.

Which of the following tools can the data engineer use to solve this problem?

A. Auto Loader

B. Unity Catalog

C. Delta Lake

D. Delta Live Tables

Show Suggested Answer

by 806e7d2 at *Nov. 21, 2024, 10:21 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 117 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 117
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer has three tables in a Delta Live Tables (DLT) pipeline. They have configured the pipeline to drop invalid records at each table. They notice that some data is being dropped due to quality concerns at some point in the DLT pipeline. They would like to determine at which table in their pipeline the data is being dropped.

Which approach can the data engineer take to identify the table that is dropping the records?

A. They can set up separate expectations for each table when developing their DLT pipeline.

B. They can navigate to the DLT pipeline page, click on the "Error" button, and review the present errors.

C. They can set up DLT to notify them via email when records are dropped.

D. They can navigate to the DLT pipeline page, click on each table, and view the data quality statistics.

Show Suggested Answer

by 9d4d68a at *Aug. 27, 2024, 11:21 a.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 118 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 118
Topic #: 1
[All Certified Data Engineer Associate Questions]

What is used by Spark to record the offset range of the data being processed in each trigger in order for Structured Streaming to reliably track the exact progress of the processing so that it can handle any kind of failure by restarting and/or reprocessing?

A. Checkpointing and Write-ahead Logs

B. Replayable Sources and Idempotent Sinks

C. Write-ahead Logs and Idempotent Sinks

D. Checkpointing and Idempotent Sinks

Show Suggested Answer

by 9d4d68a at *Aug. 27, 2024, 12:05 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 119 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 119
Topic #: 1
[All Certified Data Engineer Associate Questions]

What describes the relationship between Gold tables and Silver tables?

A. Gold tables are more likely to contain aggregations than Silver tables.

B. Gold tables are more likely to contain valuable data than Silver tables.

C. Gold tables are more likely to contain a less refined view of data than Silver tables.

D. Gold tables are more likely to contain truthful data than Silver tables.

Show Suggested Answer

by 9d4d68a at *Aug. 27, 2024, 12:02 p.m.*

Comments

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 120 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 120
Topic #: 1
[All Certified Data Engineer Associate Questions]

What describes when to use the CREATE STREAMING LIVE TABLE (formerly CREATE INCREMENTAL LIVE TABLE) syntax over the CREATE LIVE TABLE syntax when creating Delta Live Tables (DLT) tables using SQL?

A. CREATE STREAMING LIVE TABLE should be used when the subsequent step in the DLT pipeline is static.

B. CREATE STREAMING LIVE TABLE should be used when data needs to be processed incrementally.

C. CREATE STREAMING LIVE TABLE should be used when data needs to be processed through complicated aggregations.

D. CREATE STREAMING LIVE TABLE should be used when the previous step in the DLT pipeline is static.

Show Suggested Answer

by 9d4d68a at *Aug. 27, 2024, 12:05 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 121 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 121
Topic #: 1
[All Certified Data Engineer Associate Questions]

A Delta Live Table pipeline includes two datasets defined using STREAMING LIVE TABLE. Three datasets are defined against Delta Lake table sources using LIVE TABLE.

The table is configured to run in Production mode using the Continuous Pipeline Mode.

What is the expected outcome after clicking Start to update the pipeline assuming previously unprocessed data exists and all definitions are valid?

A. All datasets will be updated at set intervals until the pipeline is shut down. The compute resources will persist to allow for additional testing.

B. All datasets will be updated once and the pipeline will shut down. The compute resources will persist to allow for additional testing.

C. All datasets will be updated at set intervals until the pipeline is shut down. The compute resources will be deployed for the update and terminated when the pipeline is stopped.

D. All datasets will be updated once and the pipeline will shut down. The compute resources will be terminated.

Show Suggested Answer

by 9d4d68a at *Aug. 27, 2024, 12:15 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 122 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 122
Topic #: 1
[All Certified Data Engineer Associate Questions]

Which type of workloads are compatible with Auto Loader?

A. Streaming workloads

B. Machine learning workloads

C. Serverless workloads

D. Batch workloads

Show Suggested Answer

by RandomForest at *Oct. 18, 2024, 11:16 a.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 123 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 123
Topic #: 1
[All Certified Data Engineer Associate Questions]

A data engineer has developed a data pipeline to ingest data from a JSON source using Auto Loader, but the engineer has not provided any type inference or schema hints in their pipeline. Upon reviewing the data, the data engineer has noticed that all of the columns in the target table are of the string type despite some of the fields only including float or boolean values.

Why has Auto Loader inferred all of the columns to be of the string type?

A. Auto Loader cannot infer the schema of ingested data

B. JSON data is a text-based format

C. Auto Loader only works with string data

D. All of the fields had at least one null value

Show Suggested Answer

by csrazdan at *July 19, 2024, 4:59 a.m.*

Comments

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 124 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 124
Topic #: 1
[All Certified Data Engineer Associate Questions]

Which statement regarding the relationship between Silver tables and Bronze tables is always true?

A. Silver tables contain a less refined, less clean view of data than Bronze data.

B. Silver tables contain aggregates while Bronze data is unaggregated.

C. Silver tables contain more data than Bronze tables.

D. Silver tables contain less data than Bronze tables.

Show Suggested Answer

by helmerpaiva at *May 10, 2024, 1:31 a.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 125 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 125
Topic #: 1
[All Certified Data Engineer Associate Questions]

Which query is performing a streaming hop from raw data to a Bronze table?

A.
```
(spark.table("sales")
    .groupBy("store")
    .agg(sum("sales"))
    .writeStream
    .option("checkpointLocation", checkpointPath)
    .outputMode("complete")
    .table("newSales")
)
```

B.
```
(spark.read.load(rawSalesLocation)
    .write
    .mode("append")
    .table("newSales")
)
```

C.
```
(spark.table("sales")
    .withColumn("avgPrice", col("sales") / col("units"))
    .writeStream
    .option("checkpointLocation", checkpointPath)
    .outputMode("append")
    .table("newSales")
)
```

D.
```
(spark.readStream.load(rawSalesLocation)
    .writeStream
    .option("checkpointLocation", checkpointPath)
    .outputMode("append")
    .table("newSales")
)
```

Show Suggested Answer

by CommanderBigMac at *Sept. 12, 2024, 6:08 p.m.*

EXAM CERTIFIED DATA ENGINEER ASSOCIATE TOPIC 1 QUESTION 126 DISCUSSION

Actual exam question from Databricks's Certified Data Engineer Associate
Question #: 126
Topic #: 1
[All Certified Data Engineer Associate Questions]

A dataset has been defined using Delta Live Tables and includes an expectations clause:

CONSTRAINT valid_timestamp EXPECT (timestamp > '2020-01-01') ON VIOLATION DROP ROW

What is the expected behavior when a batch of data containing data that violates these constraints is processed?

A. Records that violate the expectation cause the job to fail.

B. Records that violate the expectation are added to the target dataset and flagged as invalid in a field added to the target dataset.

C. Records that violate the expectation are dropped from the target dataset and recorded as invalid in the event log.

D. Records that violate the expectation are added to the target dataset and recorded as invalid in the event log.

Show Suggested Answer

---

by PreranaC at *May 22, 2024, 10:03 p.m.*